

Chidera's Deliverables

What Chidera owns, what is finished, what is still to do, and the number that says each item is done. The endpoints Paul's dashboard depends on are in the companion document, *Dashboard API Contract*.

Chidera owns everything between a payment arriving and a person seeing it: the intake service, the scoring queue and workers, the live measurements, the rules and policy code, the case system and response clocks, the recommendation engine, sign-in and roles, the permanent record, the database, and the API the dashboard reads. Ire supplies the data and the trained model; Paul builds the screens on top of this API. Chidera's job is that every payment is scored once, fast, reproducibly, and that the screens can only do what each role is allowed to do.

Where things stand

AREA	MEASURE	TARGET	NOW	STATUS
Speed	Payment to decision, 95% of the time (NFR-001)	under 2 s	0.88 s	PASS · RE-RUN AFTER EVENTS
Load	Payments lost in a five-times burst (NFR-002)	0	0	PASS
Workload	Alerts per day, whole system	≤ 75	72 in the rebuilt demo day (15 Sep)	MET
Record	Entries that can be edited or deleted	none	none · hash-chained, checkable	MET
Measurements	Same twenty-two numbers live as in training, events and credits included	identical	identical, tested	MET
Roles	Administrator can open a case	never	never (refused by the server)	MET
Sign-in	Two-step sign-in for lead and administrator	mandatory	mandatory	MET
Advice	Every case has an action, steps and a reason	100%	100%	MET
Advice	How often analysts follow the advice	reported weekly	recorded, not reported	NOT YET

AREA	MEASURE	TARGET	NOW	STATUS
Advice	Bank's own procedure loaded	yes	our playbook, replaceable	PARTLY
Dashboard	Endpoints the full dashboard needs	all	four gaps listed	PARTLY
Governance	False alarms per rule and per model (base PRD FR-505)	reported	built; demo evidence thin	MET
Governance	Every rule has an owner, reason and review date (FR-504)	all	all nine	MET
Record	Who opened each case is recorded (FR-408)	every view	every view	MET
Quality	Automated tests, against a real database	all pass	198 pass	MET
Regulation	Refund and notification clocks in working days (REG-NG-05, -06)	built	built, breaches recorded once	MET
Regulation	24-hour BVN flag, shared across banks (REG-NG-04)	built	built to the connection; NIBSS link needs a bank	PARTLY
Workload	Analyst reviews a day after machine actions	below 75	51 (measured, D80a)	SHADOW FIRST

Finished — Chidera presents these at the defence

- **Intake.** Single and batch payment endpoints; a repeated reference is recognised as a duplicate, not a second payment. Identities are one-way hashed on arrival and never stored readable (D9c, D19a).
- **Queue and workers.** A database-backed queue with leases, so no payment is handed to two workers; a test proves it. Three workers keep up with a five-times burst without losing a payment.
- **Scoring.** Measurements are computed once, then go to the rules and the model at the same time; the policy step alone combines their answers (see the Blueprint diagram). Every decision stores the versions that produced it, so it can be reproduced months later.
- **Rules and policy.** Nine rules of three powers (raise, override, lower), each for the direction it reads: outgoing payments, incoming credits, or both. The payment-burst rule looks at where money is going (D67). The whole-system alert budget, 75 a day (D76), spends rule alerts first and gives the model what is left.

- **One event, not one transaction (D72).** An event envelope for sign-ins, device binding, credential and SIM changes, payee additions and credits, validated and tokenised at the same boundary as payments; the account takeover sequence rule reads them (D77).
- **Money coming in (D78, D79).** Credits are ingested and decided by receiving-side rules, never by the payment model; the payment that moves a fan-in on is caught as the second leg and lands in the same case.
- **Regulatory clocks (D71).** Customer report time, milestones recorded once, clocks in Lagos working days with public holidays, and a breach recorded once and escalated.
- **The 24-hour flag and BVN identity (D73, D75).** A flag that ends on its own and escalates a missed contact; BVN tokens stamped at intake through replaceable adapters for the core banking system, the identity registry and the industry watch-list, each running on a stand-in until a bank connects.
- **Directives a bank can act on (D74).** A signed, versioned decision per payment written in the scoring transaction, with fetch, feed and acknowledgement endpoints, failing open, in shadow until a bank turns it on.
- **Dispositions (D80).** Every actionable decision is marked machine action, analyst review or auto-close from a versioned policy; machine-handled cases stay out of "take next", and Operations shows the day by tier.
- **Case system.** Alerts about one customer group into one case within 24 hours; response clocks per severity; a worklist ordered by severity, money and lateness; "take the next case" assigns it so two analysts never work the same one.
- **Recommendation engine.** One action, numbered steps and a reason for every case, from a written playbook; the analyst's choice of whether they followed it is saved with every outcome.
- **Sign-in, roles and the record.** Four roles, server-side sessions in a cookie scripts cannot read, two-step sign-in, and a permanent record where every entry is chained to the last; a check endpoint proves nothing was altered.
- **Live alerts.** A stream that replays anything missed after a dropped connection, so the dashboard never loses an alert.
- **Governance from the base PRD review (D69).** False alarms per rule and per model with 95% ranges; a rules inventory with owners, reasons, review dates and flags for rules with no owner, an overdue review or no firing; median time to resolve; and a record of every case view.
- **Demo builder.** Rebuilds a believable desk in about ten minutes; outcomes follow the simulator's truth (D66a) and the live day continues the same bank (D68).

Still to do

1 Close the four API gaps the full dashboard needs AGREED WITH PAUL

Why	The shareable Risk Radar Desk design has views the current API cannot fully feed. Each gap, with its proposed endpoint, is in the <i>Dashboard API Contract</i> .
What	Advice-followed rate on the operations metrics; the model's evaluation figures for the administrator; an "escalated to" filter for information security; dormancy on worklist items. ("Turned into alerts" per rule is done.)
Done when	Each field exists, has a test, and Paul confirms the screen reads it with no workaround.

2 Report how often analysts follow the advice NOT YET

Why	It is the only honest check on whether the recommendation engine is any good. The answer is already saved with every outcome; nothing reads it back.
What	Add to GET /v1/metrics/operations: share of decided cases where the advice was followed, overall and per recommended action.
Done when	The ops-lead view shows the rate for the last seven days, and the Blueprint's "not yet" becomes a number.

3 Put the team's alert-budget choice live DONE · 75 A DAY

Why	A smaller budget cuts false alarms sharply now that the burst rule is precise: about 75 a day at 120 alerts, 33 at 75, 20 at 60, at some cost to never-seen account takeover (D67c).
What	Once chosen: set the budget, run <code>scripts/derive_thresholds.py --budget N --publish</code> for a new audited threshold set, rebuild the demo, and update the figures in every document.
Done when	The live threshold set records the new budget, the demo day stays within it, and the decision is logged. Done: D76; the demo was rebuilt on 15 September. The live day raised 72 alerts against 75 and caught all three fraud incidents in it (one mule fan-out, two card testing). The first attempt was refused by the threshold script: measured over the whole month, the receiving-side rules alone came to 532 a day while the credit window filled. Thresholds now come from the last 15 days, and a failed reset resumes with <code>--from-step --anchor</code> (D78f).

4 Make the bank's procedure loadable SLOT BUILT

Why	No bank has shared its fraud procedures, so the playbook is our own. It must be replaceable without code, and never described as a bank's policy until one is loaded.
What	Document the playbook's format (situation in, action, steps and urgency out) and load it from a file the administrator can replace, with every change recorded.
Done when	Swapping the file changes the recommendations on the next case, with no deployment, and the change appears in the permanent record.

5 Keep the contract and the code in step ONGOING

Why	A field renamed on the server breaks a screen silently.
What	A contract test that calls each endpoint in the API contract and checks the fields listed there; any change to a field updates the contract document in the same commit.
Done when	The test runs with the others, and the contract is frozen on the Monday of week five (feature freeze).

6 Prepare the Nigerian regulatory clocks BUILT · VALUES TO CONFIRM

Why	The enterprise base PRD's largest Nigerian gap: scam-refund clocks (notify the receiving bank in 30 minutes, refund within 16 working days of first report, as that PRD states them) and a BVN watchlist with a 24-hour temporary flag (D69g).
What	Confirm the clock values against the CBN circular with Kanyinsola; then design clocks as versioned configuration rather than code, with a case field for the customer's first report time.
Done when	A reviewed design with confirmed values and their source is in the decision log. Built as versioned configuration (D71), with the 24-hour flag (D73). Confirming the values against the CBN circular with Kanyinsola is still open; changing them is a new policy version, not code.

Rules Chidera works under

- **Advisory only.** Nothing Risk Radar does blocks, holds or reverses a payment. A directive or machine action is a recommendation the bank's own systems choose to act on (D74, D80c).
- **The server enforces roles.** The dashboard hides what a role cannot do, but the API refuses it regardless.
- **Money is integer kobo everywhere,** in the database, the API and every total.
- **Decisions and alerts are never edited;** only cases change, and every change is recorded.
- **Detection never reads the simulator,** and a test enforces it.