

# Chidera's Deliverables

What Chidera owns, what is finished, what is still to do, and the number that says each item is done. The endpoints Paul's dashboard depends on are in the companion document, *Dashboard API Contract*.

**Chidera owns everything between a payment arriving and a person seeing it:** the intake service, the scoring queue and workers, the live measurements, the rules and policy code, the case system and response clocks, the recommendation engine, sign-in and roles, the permanent record, the database, and the API the dashboard reads. Ire supplies the data and the trained model; Paul builds the screens on top of this API. Chidera's job is that every payment is scored once, fast, reproducibly, and that the screens can only do what each role is allowed to do.

## Where things stand

| AREA         | MEASURE                                                  | TARGET          | NOW                            | STATUS  |
|--------------|----------------------------------------------------------|-----------------|--------------------------------|---------|
| Speed        | Payment to decision, 95% of the time (NFR-001)           | under 2 s       | 0.88 s                         | PASS    |
| Load         | Payments lost in a five-times burst (NFR-002)            | 0               | 0                              | PASS    |
| Workload     | Alerts per day, whole system                             | ≤ 120           | 82 (latest demo day)           | MET     |
| Record       | Entries that can be edited or deleted                    | none            | none · hash-chained, checkable | MET     |
| Measurements | Same twelve numbers live as in training                  | identical       | identical, tested              | MET     |
| Roles        | Administrator can open a case                            | never           | never (refused by the server)  | MET     |
| Sign-in      | Two-step sign-in for lead and administrator              | mandatory       | mandatory                      | MET     |
| Advice       | Every case has an action, steps and a reason             | 100%            | 100%                           | MET     |
| Advice       | How often analysts follow the advice                     | reported weekly | recorded, not reported         | NOT YET |
| Advice       | Bank's own procedure loaded                              | yes             | our playbook, replaceable      | PARTLY  |
| Dashboard    | Endpoints the full dashboard needs                       | all             | four gaps listed               | PARTLY  |
| Governance   | False alarms per rule and per model (base PRD FR-505)    | reported        | built; demo evidence thin      | MET     |
| Governance   | Every rule has an owner, reason and review date (FR-504) | all             | all six                        | MET     |

| AREA    | MEASURE                                   | TARGET     | NOW        | STATUS |
|---------|-------------------------------------------|------------|------------|--------|
| Record  | Who opened each case is recorded (FR-408) | every view | every view | MET    |
| Quality | Automated tests, against a real database  | all pass   | 103 pass   | MET    |

## Finished — Chidera presents these at the defence

- **Intake.** Single and batch payment endpoints; a repeated reference is recognised as a duplicate, not a second payment. Identities are one-way hashed on arrival and never stored readable (D9c, D19a).
- **Queue and workers.** A database-backed queue with leases, so no payment is handed to two workers; a test proves it. Three workers keep up with a five-times burst without losing a payment.
- **Scoring.** Measurements are computed once, then go to the rules and the model at the same time; the policy step alone combines their answers (see the Blueprint diagram). Every decision stores the versions that produced it, so it can be reproduced months later.
- **Rules and policy.** Six rules of three powers (raise, override, lower). The payment-burst rule now looks at where money is going (D67). The whole-system alert budget spends rule alerts first and gives the model what is left.
- **Case system.** Alerts about one customer group into one case within 24 hours; response clocks per severity; a worklist ordered by severity, money and lateness; "take the next case" assigns it so two analysts never work the same one.
- **Recommendation engine.** One action, numbered steps and a reason for every case, from a written playbook; the analyst's choice of whether they followed it is saved with every outcome.
- **Sign-in, roles and the record.** Four roles, server-side sessions in a cookie scripts cannot read, two-step sign-in, and a permanent record where every entry is chained to the last; a check endpoint proves nothing was altered.
- **Live alerts.** A stream that replays anything missed after a dropped connection, so the dashboard never loses an alert.
- **Governance from the base PRD review (D69).** False alarms per rule and per model with 95% ranges; a rules inventory with owners, reasons, review dates and flags for rules with no owner, an overdue review or no firing; median time to resolve; and a record of every case view.
- **Demo builder.** Rebuilds a believable desk in about ten minutes; outcomes follow the simulator's truth (D66a) and the live day continues the same bank (D68).

## Still to do

### 1 Close the four API gaps the full dashboard needs AGREED WITH PAUL

|           |                                                                                                                                                                                                                               |
|-----------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Why       | The shareable Risk Radar Desk design has views the current API cannot fully feed. Each gap, with its proposed endpoint, is in the <i>Dashboard API Contract</i> .                                                             |
| What      | Advice-followed rate on the operations metrics; the model's evaluation figures for the administrator; an "escalated to" filter for information security; dormancy on worklist items. ("Turned into alerts" per rule is done.) |
| Done when | Each field exists, has a test, and Paul confirms the screen reads it with no workaround.                                                                                                                                      |

### 2 Report how often analysts follow the advice NOT YET

|           |                                                                                                                                                      |
|-----------|------------------------------------------------------------------------------------------------------------------------------------------------------|
| Why       | It is the only honest check on whether the recommendation engine is any good. The answer is already saved with every outcome; nothing reads it back. |
| What      | Add to GET /v1/metrics/operations: share of decided cases where the advice was followed, overall and per recommended action.                         |
| Done when | The ops-lead view shows the rate for the last seven days, and the Blueprint's "not yet" becomes a number.                                            |

### 3 Put the team's alert-budget choice live WAITING ON THE TEAM

|           |                                                                                                                                                                                              |
|-----------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Why       | A smaller budget cuts false alarms sharply now that the burst rule is precise: about 75 a day at 120 alerts, 33 at 75, 20 at 60, at some cost to never-seen account takeover (D67c).         |
| What      | Once chosen: set the budget, run <code>scripts/derive_thresholds.py --budget N --publish</code> for a new audited threshold set, rebuild the demo, and update the figures in every document. |
| Done when | The live threshold set records the new budget, the demo day stays within it, and the decision is logged.                                                                                     |

### 4 Make the bank's procedure loadable SLOT BUILT

|           |                                                                                                                                                                       |
|-----------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Why       | No bank has shared its fraud procedures, so the playbook is our own. It must be replaceable without code, and never described as a bank's policy until one is loaded. |
| What      | Document the playbook's format (situation in, action, steps and urgency out) and load it from a file the administrator can replace, with every change recorded.       |
| Done when | Swapping the file changes the recommendations on the next case, with no deployment, and the change appears in the permanent record.                                   |

## 5 Keep the contract and the code in step ONGOING

|           |                                                                                                                                                                          |
|-----------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Why       | A field renamed on the server breaks a screen silently.                                                                                                                  |
| What      | A contract test that calls each endpoint in the API contract and checks the fields listed there; any change to a field updates the contract document in the same commit. |
| Done when | The test runs with the others, and the contract is frozen on the Monday of week five (feature freeze).                                                                   |

## 6 Prepare the Nigerian regulatory clocks DESIGN ONLY

|           |                                                                                                                                                                                                                                                |
|-----------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Why       | The enterprise base PRD's largest Nigerian gap: scam-refund clocks (notify the receiving bank in 30 minutes, refund within 16 working days of first report, as that PRD states them) and a BVN watchlist with a 24-hour temporary flag (D69g). |
| What      | Confirm the clock values against the CBN circular with Kanyinsola; then design clocks as versioned configuration rather than code, with a case field for the customer's first report time.                                                     |
| Done when | A reviewed design with confirmed values and their source is in the decision log. Building it waits for scope agreement.                                                                                                                        |

## Rules Chidera works under

- **Advisory only.** Nothing Risk Radar does blocks, holds or reverses a payment.
- **The server enforces roles.** The dashboard hides what a role cannot do, but the API refuses it regardless.
- **Money is integer kobo everywhere,** in the database, the API and every total.
- **Decisions and alerts are never edited;** only cases change, and every change is recorded.
- **Detection never reads the simulator,** and a test enforces it.

---

Sources: docs/qa-report.md (NFR-001, NFR-002), backend/riskradar/api/, backend/riskradar/security/rbac.py, decisions D9–D68 in docs/decisions.md. Companions: Dashboard API Contract, Ire's Deliverables, Risk Radar Blueprint, base-prd-alignment.md.