

Ire's Deliverables

What Ire owns, what is finished, what is still to do, and the number that says each item is done.

Ire owns the data foundation and the evidence for detection: the simulator that generates the bank's customers and its fraud, the twenty-six measurements taken on every payment, the sign-in and credit events behind them, the trained model, and every evaluation that says how well the rules and the model work. Chidera wires these into the live system; Ire's job is to make sure what goes in is realistic and what comes out is measured honestly.

Where things stand

AREA	MEASURE	TARGET	NOW	STATUS
Simulator	Fraud rate against Nigeria's interbank settlement body (NIBSS)	≈ 0.30%	0.279%	MET
Simulator	Fraud incidents available to test on	≥ 600	1,455 · seven kinds	MET
Simulator	Measurements that give fraud away on their own	0 over 15%	0 · worst 5.6% (first twelve)	RE-RUN FOR THE TEN NEW
Simulator	Normal behaviour checked against published Nigerian figures	key shares	fraud rate only	PARTLY
Measurements	Same twenty-six numbers in training and live (spec 1.4.1, events and credits included)	identical	identical, tested	MET
Model	Honesty of its probability (stated ÷ actual, where alerts are raised)	≥ 0.90	0.919 (15 Sep corpus)	MET
Model	Calibration error where alerts are raised	≤ 0.04	0.032	MET
Evaluation	Account takeover caught, never seen in training	≥ 0.70	1.000	MET · PARTLY WRITTEN IN (D77D)

AREA	MEASURE	TARGET	NOW	STATUS
Evaluation	Mule fan-out caught, never seen in training	—	0.960	NO FALL WITH FOUR TYPES ADDED
Evaluation	Card testing caught, never seen in training	≥ 0.90	1.000	MET
Evaluation	SIM swap · dormant account caught, never seen in training	—	0.972 · 0.872	PARTLY BY CONSTRUCTION
Evaluation	Scams · cloned cards caught, never seen in training	—	0.516 · 0.432	THE HARD CASES
Validation	Data we did not make: PaySim, retrained there, ranking vs random	better than random	101.6× (as shipped 1.9×)	MET
Validation	Two Nigerian transaction files (5.09m, 1m rows)	measurable	labels unrelated to the data (1.00–1.02× random)	UNUSABLE FOR DETECTION
Evaluation	False alerts per fraud incident, time-ordered test at 75 a day	below 12–14	2.0 : 1	MET
Evaluation	Share of fraud value detected	60–70%	82.3%	MET
Evaluation	Analyst reviews a day after machine actions, detection unchanged	below 75	57.9	SHADOW FIRST (D89C)
Evaluation	Every rule measured on its own	all thirteen	ten	PARTLY
Validation	Public dataset with invented labels (PaySim)	run	done	MET
Validation	Public dataset with real fraud labels (IEEE-CIS)	run	blocked	BLOCKED

"Never seen in training" means that kind of fraud was hidden from the model while it learned, with the team held to 75 alerts a day (D76). Average across the seven: 0.822 (decision D82b); across the first three before the new types were added it was 0.960 (D79a), and those three did not fall. On its own the model scores 0.627 across the seven; it is useful alongside the rules, not instead of them, and must always be described that way. The takeover figure is partly shaped by the simulator, which writes the sign-in events the new rule reads (D77d).

Finished — Ire presents these at the defence

- **The simulator.** Customers of four kinds (ordinary, salary earner, small business, trader) who own one to three accounts, keep habits, pay known and new people. Three kinds of fraud generated as criminal processes, not as rule thresholds: account takeover, mule fan-out, card testing. Identities now come from the seed, so the same seed rebuilds the same bank (D68).
- **Four Nigerian fraud types (D82):** scams paid into a collection account, SIM swap drains, dormant accounts, cloned cards, each with a legitimate twin (group collections, SIM replacements, owners returning, travel), four new measurements, and rule settings chosen by a stated criterion. The tuning caught three faults before anything shipped.
- **The dataset evaluator (D81):** any transaction file mapped and measured through every layer at the budget, with the unseen-type test and a capability report; run on PaySim and two Nigerian files, with a label diagnostic proving those files unusable for detection. IEEE-CIS is mapped and ready.
- **The twenty-six measurements,** one shared package for training and live, with a test proving they match. Version 1.4.1 adds region novelty, card-present counts, other senders to a destination and the local hour to version 1.3.0: five read sign-in, device, SIM and payee events (D77), five read money coming in (D78); "does not apply" is marked -1, ratios are capped at 100 (D64). The model reads seventeen: giving it the five receiving-side measurements cut unseen takeover from 0.991 to 0.857, so they stay with the rules (D78b).
- **The event and credit layers** in the simulator: sign-in failures, device binding, credential and SIM changes, payee additions, and inbound credits with mule rings, each on its own random stream so the payments stay byte-identical (D77c, D78d). The corpus is now anchored to a date, so it rebuilds exactly.
- **The three new detectors, each tuned and measured before going live:** account takeover sequence (D77b), inbound fan-in on a month of history (D78c), and the second leg, which took mule-ring value flagged from 63.5% to 84.9% (D79a).
- **The budget menu and the tier study:** detection, false alerts and fraud value at every budget on a time-ordered test (D70a), and the measurement that machine actions take 24 alerts a day off the analysts at precision 0.992 with detection unchanged; refilling and auto-close were tested and rejected (D80a, D80b).
- **The calibrated model:** gradient-boosted trees, trained on older data and adjusted on newer data it never saw, so its probability means what it says (D60–D61).
- **The held-out evaluation:** six approaches compared on identical data at the same alert budget, with 95% ranges on every figure (D59, D63, D65, D67a).
- **The self-audit** that caught eight "giveaway" measurements in our own data and led to the realistic population (D57), and the reproduction of the earlier dataset's result (D56).
- **PaySim validation:** 65% of fraud value caught while catching 18% of fraud transactions; six of the original twelve measurements cannot fire on PaySim, and none of the ten event and credit measurements can, which we state rather than hide.

Still to do

1 Check "normal" against published Nigerian figures NOT STARTED

Why	The most serious objection to the project is that our normal customers are invented. This turns "we made it up" into "we made it up and checked its shape".
What	Take published aggregates (NIBSS annual fraud reports, Central Bank payment statistics) and compare the simulator's shares side by side: channel mix, average transfer value, interbank (NIP) versus in-bank share, fraud loss by channel.
Done when	A table of our share against the published share for each; any gap wider than a third of the published value is either corrected in the simulator or explained in writing.

2 Run the real-label validation (IEEE-CIS) MAPPED · WAITING FOR THE FILES

Why	The only test short of a real bank with genuine fraud outcomes (about 590,000 transactions, 3.5% fraud).
What	Sign in to Kaggle, complete phone verification, accept the competition rules at kaggle.com/competitions/ieee-fraud-detection/rules , put <code>train_transaction.csv</code> in <code>ml/data/public/</code> , then run <code>python ml/validate_public.py --dataset ieee</code> . The script is written and ready.
Done when	Results file written; precision-recall area, calibration error and recall at the alert budget reported, with the caveat that it checks method, not our measurements (its columns are anonymised).

3 Measure the last three rules one at a time 10 OF 13 DONE

Why	The Blueprint requires every rule to be tested on its own and dropped if the data does not support it. Card testing (right 99.8% of the time), the payment-burst rule (D67), the familiar-payee rule (D64), and the three added since (D77b, D78c, D79) are done.
What	For "saved payee" (lowers risk): how often it fires on fraud, which would hide it. For "sanctioned destination" and "known mule destination" (override everything): seed test lists and confirm they fire on every listed destination and nowhere else.
Done when	Each of the three has a firing rate, a share of fraud it wrongly lowers or correctly catches, and a keep-or-drop recommendation logged in the decision log.

4 Give the demo bank a longer memory DONE • BURST RULE 37/DAY

Why	The payment-burst rule now fires only on payments to a destination new to the bank. The demo bank remembers about 3 days, so a third of its payments look "new" and the rule fires about 50 times a day, 5 of them on fraud; in the 30-day training data it fires about 34 times a day (D67d).
What	Lengthen the demo profile's history (in <code>simulator/riskradar_sim/generate.py</code>) from 4 days towards 21–30, and measure how long the demo rebuild then takes. Done: 30 days (D76c) . Rebuilt on 15 September: 414,239 history transactions (payments and credits) took 1 h 47 min to load and 34 min to score. In the live day the burst rule raised 37 alerts, just over the 35 aimed for; the whole system raised 72 against 75 and caught all three incidents. The receiving-side rules need the history's second half to settle (D78f).
Done when	On a rebuilt demo the rule fires at most about 35 times a day, the whole system stays within its alert budget (now 75 a day), and the rebuild time is written in the README.

5 Support the team's alert-budget decision DONE • 75 A DAY

Why	With the new rule, a smaller budget cuts false alarms sharply: 120 a day gives about 75 false alarms, 75 a day about 33, 60 a day about 20. The cost is account takeover never seen before: 0.777 at 120, 0.703 at 75, 0.658 at 60 (D67c).
What	Walk the team through the trade-off; once they choose, re-run <code>ml/evaluate_system.py</code> at the chosen budget and confirm the figures quoted in the Blueprint.
Done when	The team's choice is logged as a decision and every document quotes the figures for that budget. Done: 75 a day (D76); the held-out figures above are at 75.

Rules I re works under

- **Detection never reads the simulator.** Rules and model see only the payment and its measurements, never a generator setting; a test enforces this.
- **Publish the held-out number,** never the flattering one from fraud the model has already seen.
- **Count incidents, not transactions,** and never report accuracy: at 0.3% fraud, "never fraud" scores 99.7%.
- **Never train on someone else's column** named risk, score, rate or fraud unless the code that made it can be read (D10f).
- **Log every change** to what the numbers mean in the decision log, with its evidence, before quoting it.

Sources: `ml/artifacts/evaluation-20260915.1334-final.json`, `own-corpus-audit.json`, `public-validation-paysim.json`, `system-evaluation-75.json`, `budget-menu.json`, `tiers.json`; docs/`validation-plan.md`; decisions D56–D84 in docs/`decisions.md`. Companion: Risk Radar Blueprint.